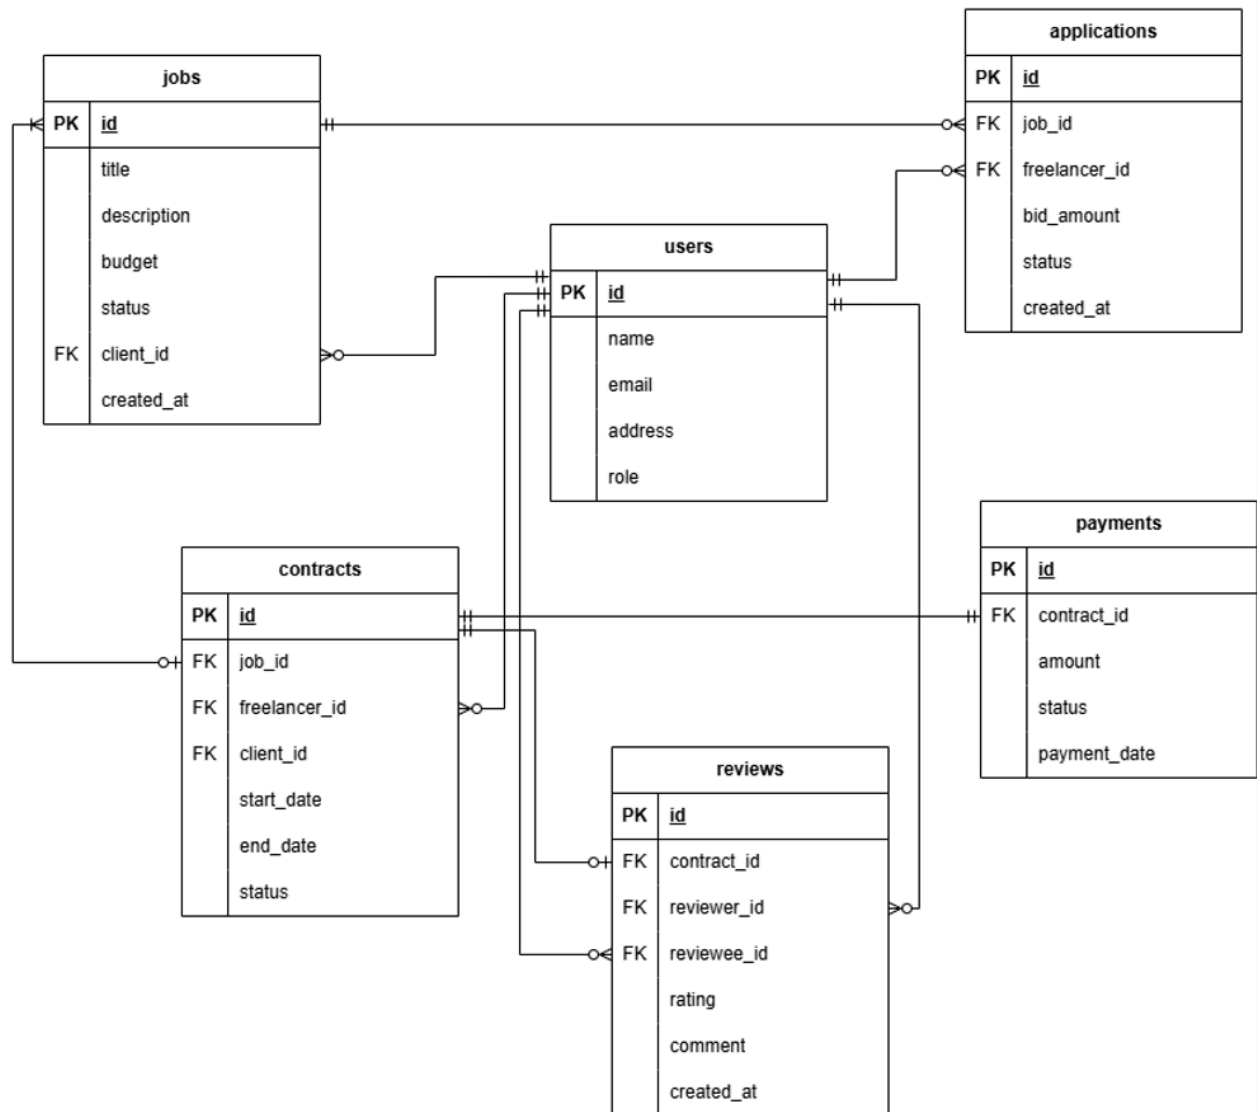


# Entity Relationship Diagram and Explanation



# Chosen Database System: PostgreSQL

For this project, I have chosen **PostgreSQL**, an advanced open-source **relational database management system (RDBMS)** known for its strong reliability, extensibility, and standards compliance. PostgreSQL is an excellent choice for a **Freelancer Job Board API** because it provides robust support for complex queries, data integrity, and scalability.

## Why PostgreSQL?

### 1. Relational Structure & ACID Compliance:

PostgreSQL is an **ACID-compliant** (Atomicity, Consistency, Isolation, Durability) database, ensuring that transactions are executed reliably. This is crucial for a job board where data integrity is essential—for example, ensuring that jobs, applications, and payments are correctly processed without data corruption.

### 2. SQL Compliance & Extensibility:

PostgreSQL follows the **SQL standard** while offering powerful extensions like **PostGIS (for geolocation)** and **JSONB** (for semi-structured data). It allows a mix of structured and unstructured data, making it flexible.

### 3. Scalability & Performance:

- PostgreSQL handles large-scale applications efficiently with indexing, partitioning, and query optimisation.
- It supports **horizontal and vertical scaling**, making it ideal for a growing application with increasing freelancer-client interactions.

### 4. Strong Security Features:

PostgreSQL provides **role-based access control (RBAC)**, SSL encryption, and built-in **authentication methods** (e.g., password-based, certificate-based, or external authentication).

### 5. ORM Compatibility (SQLAlchemy):

Since I am using **SQLAlchemy**, PostgreSQL integrates seamlessly with it, allowing me to interact with the database in an efficient and structured way.

---

# Comparison with Other Database Systems

To justify the choice of PostgreSQL, let's compare it with other popular database types:

## 1. PostgreSQL vs MySQL (Relational DB)

Feature	PostgreSQL	MySQL
ACID Compliance	Fully ACID-compliant	ACID-compliant but with some exceptions
Complex Queries	Handles complex queries well	Optimized for read-heavy workloads
JSON Support	JSON & JSONB for semi-structured data	Basic JSON support (less efficient)
Extensibility	Supports procedural languages, full-text search, and indexing	Limited extensibility
Concurrency Control	Uses MVCC (better concurrency & performance)	Uses table-level locking more often

**Verdict:** PostgreSQL is superior for complex queries, extensibility, and JSON data handling. MySQL is simpler but less feature-rich.

---

## 2. PostgreSQL vs MongoDB (NoSQL)

Feature	PostgreSQL	MongoDB
Data Structure	Relational (tables & rows)	Document-based (JSON-like BSON)
Schema Enforcement	Strongly enforced schema (structured data)	Schema-less (flexible data)
Query Language	SQL	MongoDB Query Language (MQL)
Transactions	ACID transactions	Multi-document transactions (added later)
Scalability	Vertical & horizontal scaling	High horizontal scaling

**Verdict:** If my project required highly dynamic or semi-structured data (e.g., social media feeds), MongoDB would be useful. However, since I need **structured relationships (Users, Jobs, Applications, Payments)** with strong data integrity, PostgreSQL is the better choice.

---

### 3. PostgreSQL vs Firebase (Realtime NoSQL)

Feature	PostgreSQL	Firebase
Type	Relational SQL Database	NoSQL Realtime Database
Schema	Structured & strict	Unstructured & flexible
Query Power	Advanced SQL & indexing	Basic querying
Performance	Optimized for structured data	Optimized for real-time syncing
Offline Mode	Not natively designed for offline mode	Designed for mobile apps with offline syncing

**Verdict:** Firebase is ideal for **real-time applications like messaging** but not suitable for a structured job board that requires **complex queries, relationships, and transactions**.

---

### Final Justification for PostgreSQL

Considering the project requirements—structured job listings, user accounts, applications, and financial transactions—**PostgreSQL is the best fit** because:

- It ensures **data integrity** and **supports relationships** between entities.
- It allows **complex queries** for filtering jobs, applications, and transactions.
- It scales well with indexing and optimization features.
- It integrates well with **SQLAlchemy**, making ORM-based database management easier.